2048 Expectimax Solver

Anders Sildnes, Andrej Leitner *students at NTNU*

Abstract—**This text answers assignment 5: writing a solver for the game 2048. We present a solver using Expectimax algorithm. We explain our choice of heuristics and results.**

THE purpose of "2048" is to slide tiles with numbers 2^n, and join them together to form 2048 (2^{11}), or possibly higher. Between each slide, i.e. moving all the tiles on the board north, west, east or south, a random tile valued either 2 or 4 appears on the board, which adds difficulty in solving the game.

Assumptions

It is possible to obtain tiles values up to 2^{14} within reasonable time. Our assignment bounds, however, did not require us to achieve such a high level of performance:

1) The highest value tile does not need to exceed 4096 (2^{12}).
2) Execution time is not critical (acceptable as long as < 10 minutes)

This implies that our chosen heuristic and solver method does not need a strong level of optimization to work. Indeed, the work of [1] shows that a naïve search can be used to solve the game within the bounds listed above. Still, we have implemented an expectimax approach which is much more scalable.

Minimax Algorithm

The purpose of minimax is to minimize loss for a worst possible scenario. This means choosing an action such that your opponent can do "least amount of damage", i.e. leaving you in a promising state. The process of to determine the best action available given your opponent's possible re-action is also referred to as *adversarial search*. For 2048 this would mean always making a move such that no matter where the next tile goes, and no matter what value, your action minimizes loss.

2048, however, is a stochastic game. The opponent will never intelligently select a bad location or value. Therefore, we felt that minimax, which always prepares for the worst case, is too pessimistic in its choice of actions. Furthermore, the distribution of tiles are given: there is a 90% chance of getting a tile valued 2, and 10% chance of getting a tile valued 4. Thus, it is possible to make somewhat accurate predictions about the expected likelihood of future states of the game. What cannot be predicted is the location of new tiles. However, the possible locations are typically not many enough so that an adversarial search cannot completely explore the possible outcomes.

Minimax with probabilistic values to each state leads to the modified *expectimax* algorithm. Search space is expanded similar to in minimax, but it is extended with *chance nodes*. The chance nodes considers each possible successor states, but instead of considering the value from each board/objective function, it also takes into account the probability of getting each state. Their importance is that if an action might lead to a bad, but unlikely result, this result will not significantly reduce the overall value of the action.

Heuristic

To assess whether or not a current game state is good or not, we have one heuristic function that assesses how promising a given board is. Our heuristic is based on ideas we found online and through some play of our own.

Gradients

If a large value is in the middle of the board, it will be hard to pair up with other adjacent tiles. Therefore, we would like to set a natural preference for having larger values clustered in a corner. This way, they do not interfere with other lower-valued tiles. However, having large values in separate corners

implies that there are lower valued tiles between them. This makes the board hard to maneuver and reduces the possible spaces. Therefore, we want to penalize the score when values are spread out.

To calculate the value of a board, we apply a 4×4 mask and retrieve the sum of adding all the values in the masked board. The mask is shown Figure 1. Each number in the tiles represents a scalar which is multiplied with the value from the board. Since large value might be clustered in either 4 corners, we apply 4 different masks, one for each corner respectively. Then, we only consider the largest resulting value is chosen as the one to be used as a value for the state. The procedure is described formally below:

Let

$\text{board}_{i,j} = 4 \times 4$ square lattice

s.t. $\text{board}_{0,0}$ represents north west

$dec(x) = \{x, x - 1, x - 2, x - 3\}$

$inc(x) = \{x, x + 1, x + 2, x + 3\}$

$g_{NW} = dec(3), \dots, dec(0) \ (Figure \ 1)$

$g_{SW} = dec(0), \dots, dec(3)$

$g_{SE} = inc(-3), \dots, inc(0)$

$g_{NE} = inc(0), \dots, inc(-3)$

$G = \{g_{NW}, g_{NE}, g_{SW}, g_{SE}\}$

Then:

For a given board:

$$\text{gradient} = \max_{\forall g \in G} \sum_{i,j} g_{i,j} \cdot \text{board}_{i,j}$$

3	2	1	0
2	1	0	-1
1	0	-1	-2
0	-1	-2	-3

Fig. 1: Example gradient mask (g_{NW})

Further, to each chosen gradient value, we calculate $E[\text{board}]$, i.e. the expected value, given the value of the board and the probability. This is given by the probability of having a tile valued either 2 or 4. We assume there is a uniform probability of tile location (so this probability is 1). The values are multiplied together, and when combined make up the heuristic assessment value for each board.

WHY NOT USE MORE HEURISTICS?

Many other heuristic functions exist, for example:

1) **Board score** – whenever two cells merge, the board score increases.
2) **Free tiles** – giving scores based on how many cells are left.
3) **Non-monotonic rows/columns** – prefer boards where values next to each other are in order.
4) **Possible merges** – prefer boards where many merges are possible, i.e. boards with tiles of same value.

Heuristic 1 is difficult to implement in practice. While the overall goal of the game is to achieve a high score, relying on the game score alone is a greedy choice. Sometimes it is better to keep tiles with similar values. Arguably, heuristic 2 is already used in our gradient function by adding negative values in our gradient masks. The two last heuristics are shown to be effective, but we chose not to implement them. This is for the reasons outlined below and since they were not necessary to achieve our assignment bounds.

Expectimax has execution time $O(b^m n^m)$. Here, b is the *branching factor*, i.e. the number of possible children from a given board. For each board, $0 \leq b \leq 4$. m is the depth of the search tree. n is the number of possible moves, i.e. possible tiles with value either 2 or 4. It can easily be seen that this execution time quickly grows: for a search depth of 6, the execution time is $4^6 \cdot n^4$ for each possible move. In addition, we have to apply our mask in Figure 1 four times and introduce chance nodes. Thus, to avoid cluttered code, we chose to only rely on the gradient method.

Structure of source code

We implemented our solution in Java. This is because both members of the group were familiar with the language. The GUI is rendered through a JFrame, and the tiles are drawn using AWT's Graphics2D.

There is a built-in key-listener that reads for keyboard input. However, one can also enable auto-solving, which toggles a flag in the key-listener and forces the computer to use AI to solve the game by itself. The built-in solver works as follows:

1) depth = 6
2) For each possible move (left,right,up,down) by the player:
   a) For each possible tile placement:
      i) place a tile
      ii) call all possible player moves, decrementing depth by 1
3) calculate E[current state]
4) ... goto step 1 if depth > 0
5) Get E[current node], multiply upward to root:
   a) For each maximizer node:
   b) set E[node] /= $\sum\limits_{child} p_c$
6) Once complete k-depth tree is complete, choose most promising branch.

To realize the above code, a few structures are needed. We implemented each possible board as a class with each tile as a subclass. It could be noted that since each tile will not need to exceed $2^{16} = 65536$, a 64-bit architecture can represent one row as a 64-bit number (16×4) which permits fast bit-wise operations for altering the board. However, without classes (objects) the GUI is more difficult to render, so we chose not to use this approach, even though it allows fast board checks.

Our choice of depth 6 for the search tree comes from trial and error experimenting. Depth 6 appears to have the best result to time ratio. If our search space was faster to calculate, we could probably have expanded. Conversely, if we needed higher score results, it would need to be expanded.


Fig. 2: Example expectimax tree, with implicit chance nodes. Each black cell is a board after a move is made (waiting for opponent to place a tile). The black cells' value is the highest value of its children.

Results

Our highest result achieved so far is 8192. On average, we achieve 4096 in 80% of the cases, in < 5 minutes. Figure 3 shows this result and our board.

References

[1] "ronzil" *2048-AI*, *https://github.com/ronzil/2048-AI*. GitHub

Fig. 3: Example solver and board