2048 Expectimax Solver

Anders Sildnes, Andrej Leitner *students at NTNU*

Abstract—**This text answers assignment 5: writing a solver for the game 2048. We present a solver using Expectimax algorithm. We explain our choice of heuristics and results.**

THE purpose of the game is to slide tiles with numbers 2^n, and join them together to form 2048 (2^{11}), or possibly higher. It can be thought of as a 2-player, turn-based game. The opponent places tiles valued either 2 or 4 in an availible location. Then, the other player makes a move to slide all tiles in a given direction. This goes back and forth until there are no more availible moves.

In a turn-based game, you have the time to consider the consequences for each possible action. This gives computers immense advantages over humans: IBM's chess-solving machine "Deep Blue" won against Garry Kasparov in 1997 considered more than 200 million possible moves per second[1]. In the case of 2048, this could yield a valid solution in a short amount of time. However, not everyone has access to such fast hardware and multi-threading so a way to prune the search space is needed.

Assumptions

It is possible to obtain tiles values up to 2^{14} within reasonable time. Our assignment bounds, however, did not require us to achieve such a high level of performance:

1) The highest value tile does not need to exceed 4096.
2) Execution time is not critical (roughly < 10minutes)

This implies that our chosen heuristic and solver method does not need a strong level of optimization to work. Indeed, the work of (1) shows that a naïve brute force can be used to solve the game within the bounds listed above.

Minimax Algorithm

The purpose of minimax is to minimize loss for a worst possible scenario. This means choosing an action such that your opponent can do "least amount of damages", i.e. leaving you in a promising state. This is also called *adversial search*. For 2048 this would mean always making a move such that no matter where the next tile goes, and no matter what value, you will still find a way for success.

2048, however, is a stochastic game. The opponent will never purposly select a bad tile, nor *choose* a bad value. This is to say that there is no opponent, so it is possible to always land a best case scenario. Therefore, we felt that minimax, considering the worst case, is too pessimistic in its search space. Furthermore, the distribution of numbered tiles are given: there is a 90% chance of getting a tile valued 2, and 10% chance of getting a tile valued 4. The location is random. Also note that in some cases, a tile in position i, j is no different from having a tile in position $i + i, j$ if the next move it to slide the tiles downward. Thus the possible search space is not too big and one can use statistics to get accurate enough resuls.

Using stochastic information in adversial search leads to expectimax. After expanding the search space up to a given depth, you can back up from each node and introduce a chance node. The chance nodes multiply the value to each of its children (given a state, the objective value to each of the children). The chance nodes are used when pruning the search space. Their importance is that unlikely values will yield a lower penalty to the overall score of the search node, such that the node is not necessarily pruned away.

Heuristic

To assess whether or not a current game state is good or not, we have different objective functions to assess game state. Using these to advance in the solution space is called heuristical search. Our heuristics are developed based on ideas we found online and through some play of our own. BENCHMARK BENCHMARK BENCHMARK BENCHMARK BENCHMARK BENCHMARK BENCHMARK BENCHMARK

[1]Info taken from lecture slides

Gradients

If a large value is in the middle of the board, it will be hard to pair up with other adjacent tiles. Therefore, we would like to set a natural preference for having larger values clustered in a corner. This way, they do not interfere with other lower-valued tiles. Next to the highest value tile, you want other, lower valued cells. This extends throughout the board. However, having large values in two corners implies that we have two (or more) tiles with large values, separated by lower tiles. Therefore, we penalize boards where large values are far apart.

To calculate the value of a board, we apply a 4×4 mask and retrieve the sum of adding all the values in the masked board. The mask is shown Figure 1. Each number in the tiles represents a scalar which is multiplied with the value from the board. Since large value might be clustered in all corners, we apply 4 masks, one for each corner respectively. Then, the largest resulting value is chosen as the one to be used as a value for the state.

Formally, the gradient calculation can be calculated as follows:

$$
\begin{aligned}
&\textbf{Let}\\
&\text{board}_{i,j} = 4 \times 4 \text{ square lattice}\\
&\qquad \text{s.t. board}_{0,0} \text{ represents north west}\\
&dec(x) = \{x, x-1, x-2, x-3\}\\
&inc(x) = \{x, x+1, x+2, x+3\}\\
&g_{NW} = dec(3), \ldots dec(0)\\
&g_{NE} = dec(0), \ldots dec(3)\\
&g_{SW} = inc(-3), \ldots inc(0)\\
&g_{SE} = inc(0), \ldots inc(-3)\\
&G = \{g_{NW}, g_{NE}, g_{SW}, g_{SE}\}\\
&\textbf{Then:}\\
&\textit{For a given board:}\\
&\text{gradient } = \operatorname*{argmax}_{\forall g \in G} \sum_{i,j} g_{i,j} \cdot \text{board}_{i,j}
\end{aligned}
$$

Furthermore, to each chosen gradient value, we calculate $E[\text{board}]$, i.e. the expected value, given the value of the board and the probability. This is given by the probability of having a tile valued either 2 or 4. We assume there is a uniform probability of tile location (so this probability is 1). The values are multiplied together, and when combined make up the heuristical assessment value for each board.

WHY NOT USE MORE HEURISTICS?

Many other heuristic functions exist:
1) **Board score** – whenever two cells merge, the board score increases.
2) **Free tiles** – giving scores based on how many cells are left.
3) **Non-monotnic rows/coumns** – prefer boards where there are large, but similar tiles.
4) **Possible merges** – prefer boards where many merges are possible, i.e. boards with tiles of the same value.

Heuristic 1 is difficult to implement in practise. When greedily preffering boards with higher tiles you might exclude boards with high merges. Heuristic 2 can also be used. Arguably, it is already used in our gradient function: boards with tiles in opposite corners are penalized. Therefore we did not see the need to implement this heuristic expliclitly. Furthermore, there are cases where having many free tiles not necessarily is a bad thing. The two last heuristics are shown to

3	2	1	0
2	1	0	-1
1	0	-1	-2
0	-1	-2	-3

Fig. 1: Example gradient mask (g_{NW})

Exepctimax has execution time $O(b^m n^m)$. Here, b is the *branching factor*, i.e. the number of possible children from a given board. For each state, $0 \leq b \leq 4$. m is the depth of the search tree. n is the number of possible moves, i.e. possible tiles with value either 2 or 4. It can easily be seen that this execution time quickly grows: for a search depth of 6, the execution time is $4^6 \cdot n^4$ for each possible move. In addition, we have to apply our mask in Figure 1 four times and introduce chance nodes. Therefore, to avoid cluttered code and execution, we chose to only rely on the gradient method. However, if this assignment had rewarded execution time and scores higher than 4096, we would implement more of the heuristics mentioned above.

Structure of source code

We implemented our solution in Java. This is because both members of the group were familiar with the language. The GUI is rendered through a JFrame, and the tiles are drawn using AWT's Graphics2D.

There is a builtin keylistener that reads for keyboard input. However, one can also enable autosolving, which toggles a flag in the keylistener and solves the board by itself. The builtin solver works as follows:

1) For each possible move (left,right,up,down) by the player, call "computers"
2) "computer": For each possible tile placement:
   a) place a tile
   b) call all possible player moves, decrementing depth by 1
3) calculate E[current state]
4) ... goto step 1 if depth > 0
5) Get E[current node], multiply upward to root
6) Once complete k-depth=tree is complete, choost most promising branch.

To realize the above code, a few classes are needed. We implemented each possible board as a class with each tile as a subclass. It could be noted that since each tile will not need to exceed $2^{16} = 65536$, a 64-bit architecture can represent one row as a 64-bit number (16×4) which permits fast bitwise operations for altering the board. However, without classes (objects) the GUI is more difficult to render, so we chose not to use this apporach, even though it allows fast board checks.


Fig. 2: Example expectimax tree, with implicit chance nodes. Each black cell is a board after a move is made (waiting for opponent to place a tile). The black cells' value is the highest value of its children.

Results

Our highest result achieved so far is 8192. On average, we achieve 4096 in 80% of the cases, in < 5 minutes. Figure 3 shows this result and our board.

References

[1] "ronzil" *2048-AI*, *git-repository*. GitHub

Fig. 3: Example solver and board